



@andrewcmorrow #nyccpp

NYC C++ Jan 13 Meetup: error_or

Andrew Morrow -- acm@10gen.com

Member Technical Staff, 10gen (the MongoDB company)

Caveats

- I'm still learning about this stuff
- I'm almost surely wrong about some things here.
- If you think I'm wrong about something, please tell me:
 - I don't want to spread misinformation
 - I learn something too

Why not throw?

- Because you work at Google.
 - Because you previously worked at Google.
 - Because your team lead used to work at Google.
- Because the framework you are using is not exception safe.
- Because you are on a constrained environment and your runtime doesn't offer it.
- Because writing exception safe code is really hard.

Error handling without exceptions

- Return an error code and use out parameters.
 - Loses return slot, awkward to call.
 - <https://gist.github.com/4589211>
- Return a value, and have a (thread local?) `getLastError` function.
 - You must return a “bad” value, easy to forget to check error.
 - <https://gist.github.com/4589420>
- Return `std::pair<error_type, value_type>`
 - Uses extra space; how do I know which is valid?
 - <https://gist.github.com/4589450>

Pretty OK: return boost::variant<T, E>

- Like returning std::pair, but only contains either T value or E value.
 - Requires dependency on boost.
 - boost::variant is complicated (non-empty guarantee).
 - We can do better by using C++11 and rolling our own type.

error_or<E, V>

- Code: https://github.com/acmorrow/error_or
 - Requires C++11
 - Final, new style unions, move semantics, noexcept, [std::error_code](#), etc.
 - Logical successor to boost::variant approach.
 - Stores either an E value to represent an error, or a V value to represent a returned value.
 - Intended to be pure “type calculus”, it should compile away to operations on the underlying types, minimal perf hit.
 - Designed for use with [std::error_code](#) or [std::error_condition](#) as the ‘E’ type, but would work with others.

Note on code links

- The links here are to a specific revision to preserve line numbers:
 - https://github.com/acmorrow/error_or/tree/f7c0092c02b153f73e8dfb806f0a3ec9018a2593
 - Subsequent revisions may have slightly different line numbers but should be in the same relative position in the file.
 - Code links in this presentation will list the file and line if you want to view the code on github during the talk.
 - Head revision may slightly differ

Declaration

- [error_or.hpp: L28-L33](#)

- First template argument is the error type that is returned to indicate failure:
 - Just about any type will do, but intended to be `std::error_code` or `std::error_condition`.
- Second argument is the value returned on success.
- Class is final so we know there are no subclasses, no need for a virtual dtor, etc. Almost all calls should inline away.
- Add STL style typedefs for T (value_type) and E (error_type) so that metaprogramming can be done with this class.

Review: Placement new

- ‘new expression’ normally does two things:
 - Obtain a region of appropriately sized and aligned memory
 - Invokes the constructor
- Placement new:
 - Does not allocate dynamic memory.
 - Takes an argument, which is an address.
 - Returns the address it is passed.
 - Invokes the constructor at that returned address.
- Placement new allows you to construct an object into a given region of existing memory.

Review: Explicit call to destructor

- Complement to placement new:
 - Invokes the destructor directly.
 - `t->T::~~T();`
 - `t.~T();`
 - Does not invoke dynamic memory allocator.
- Allows us to 'tear down' an object that exists at a certain memory location, leaving that memory region uninitialized (and therefore able to host a new object in the future)

Members

- [error_or.hpp: L176-L185](#)
 - ‘val_’: A union of non-trivial class types (a.k.a ‘unrestricted unions’)
 - Not permitted in C++03, OK in C++11
 - Requires that we use placement new and explicit calls to the dtor to initialize and destroy members.
 - The empty val ctor/dtor inhibit the compiler from generating members for us.
 - ‘ok_’: default initialized in the class body, tells us which field of the union is populated (‘discriminator’).

Lifecycle

- C++03: [Rule of 3](#)
 - dtor, copy ctor, copy assign operator: nontrivial existence of any suggests need for all.
 - For templates, maybe include converting copy ctor and converting move assign operator.
- C++11: Rule of 5?
 - All Rule of 3 members plus new special functions: move constructor, move assign operator
 - For templates, maybe include converting move ctor and converting move assign.

C++11: Rule of Zero when possible

- Don't write lifecycle members unless you need to
 - If all of your class members have good semantics, you may not need to write any special members yourself anyway.
 - Use containers, C++11 smart pointers, RAII, etc.
 - Use '=default' if you want to be explicit.
 - Use '=delete' to disable behavior.
- For error_or though we need to be explicit since we manage memory directly (the union slots).

C++11 noexcept

- Adding 'noexcept' specifier to a function states that it will not throw.
 - Well, actually, that if it does throw, to call std::terminate.
 - Much simpler codegen than honoring exception specifiers.
 - Exception specifiers are deprecated in C++11.
- For a template it may depend on properties of template type parameters:
 - noexcept operator: `noexcept(<expression>)`
 - std::is_nothrow_[action]able

Default Constructor

- [error_or.hpp: L45-L47](#)
 - Only works if `value_type` has a default constructor
 - We use placement new to initialize the appropriate (in this case 'value' obviously) member of the union.
 - 'ok_' is true because of the in class body initializer.
 - Is noexcept, as long as `value_type`'s default constructor is noexcept.
 - The remaining constructors use `std::move` and move semantics...

What does [std::move](#) do anyway?

- It doesn't 'move' anything, it **enables** moves:
 - “obtains an rvalue reference to its argument”
 - The return value of [std::move](#) can bind to an argument that takes an rvalue reference.
 - `void some_func(Foo&& f) {`
 - `Foo wrong(f); // NO: calls copy ctor!`
 - `Foo right(std::move(f)); // YES: calls move ctor.`
 - ‘f’ is ‘dead’ at this point.
 - When in doubt: `&& => use std::move`
 - Unless you need [std::forward](#) (not covered today)
 - See Meyer’s talk on ‘[Universal References](#)’ for details

Value and Error Constructors

- [error_or.hpp: L49-L57](#)

- If the argument is an `error_type`, then we set `'ok_'` to false, and placement new into the `val_.error` field.
- If the argument is a `value_type`, then we placement new into the `val_.value` field.
- We take the arguments by value (but may be zero copy!).
- We enable moving from the copied arguments by using `std::move`.
 - The copied arg is private, so it's OK if we destroy it.

Destructor

- [error_or.hpp: L135-L141](#)
 - We switch on 'ok_' to determine which union slot is currently populated, and then make an explicit call to the destructor on the value in the indicated slot.
 - Good 'noexcept' example: this function is noexcept if the dtor's for both error_type and value_type are noexcept.
 - We can use [std::is_nothrow_destructible](#) to find out.
 - Use the result of the above in a noexcept expr.

Copy Constructor

- [error_or.hpp: L64-L71](#)
 - Argument is passed by const&, as you would expect.
 - Introspect 'other.ok_' to determine which field in other is valid, and use placement new to initialize the same field in our union with a copy of the appropriate field from other.
 - this->ok_ gets the same value as other.ok_
 - Why can't we take these arguments by value then move, like we did before?

Move Constructor

- [error_or.hpp: L73-L82](#)
 - Much like copy ctor, but we want to move from the rvalue reference.
 - Use of [std::move_if_noexcept](#) may be too conservative, using `std::move` here might be fine.
 - `std::move_if_noexcept` works like [std::move](#), but if the type's move constructor is not `noexcept`, returns `'const T&'`, inhibiting moveage.

What about assignment?

- Construction is “easy”: copy or move the provided argument into the correct union slot with placement new.
- Assignment is hard: the target error_or already has state that we must tear down.
 - We would like strong exception safety
 - ‘cis’ assignments (error valued to error valued, value valued to value valued) are relatively straightforward.
 - ‘trans’ assignments are hard
 - What if one move works, but the other throws?
 - If we catch, and move back, and *that* throws?

Use 'copy and swap' for assign

- Copy and swap is the gold standard for exception safe copy assignment operators:
 - Make a copy of the thing you are assigning from
 - It's ok, you will have two at the end anyway.
 - Swap your internals with the copy (this is cheap).
 - Copy goes out of scope (destruction was inevitable).
 - Relies on exception safety of swap.
- We have merely pushed the problem to swap.
 - Solution: Use SFINAE to disable swap, and therefore assign, if we can't build a non-throwing swap over E and T.

Copy Assignment Operators

- [error_or.hpp: L114-L127](#)
 - Take arguments by value to achieve ‘copy’ step.
 - Again, may in fact not require any copying.
 - Invoke our move constructor to steal the copy from the argument.
 - Swap our internals with the newly constructed error_or.
 - All is well, as long as swap won’t throw...

Exception safe swap?

- [error_or.hpp: L187-L210](#)
 - Maybe. We use [std::enable_if](#) to SFINAE this swap in or out depending on a class constant 'is_nothrow_swappable'
 - Requirements for error_or::is_nothrow_swappable:
 - [error_or.hpp: L36-L42](#)

Review: ADL Swap

- Old advice: specialize `std::swap` in namespace `std` for your type
 - Ugly, and it doesn't work for templates anyway.
- Better way: leverage 'argument dependent lookup' to find `swap` in your own namespace
 - Clean, and works for templates
 - Just write your in-namespace `swap` as a two argument free function.
- Relies on callers always invoking `swap` with no qualifier.
 - Say 'using `std::swap`' first so you fall back to `std` variant.

Exception safe swap (part 2)

- For 'cis' swap, we can swap with no exception if `value_type` and `error_type` can swap with no exception.
 - There is no `std::is_[nothrow_]swappable!`
 - This seems like an oversight in the standard?
 - We will roll our own:
 - [is_nothrow_swappable.hpp](#)
 - Pretty complex. Is it correct?

detail::is_[nothrow_]swappable

- Enable ADL swaps by 'using std::swap' in a private namespace
- SFINAE on decltype(swap(..., ...)) to check if swappable
 - Build an integral constant from the result
- Use noexcept(swap(..., ...)) to check for nothrow swap, but only if the underlying type is swappable.
- Cobbled together from various sources,

Exception safe swap (part 3)

- For 'trans' swap, we need noexcept move and dtor
 - The standard gives us these
 - `std::is_nothrow_move_constructible`
 - `std::is_nothrow_destructible`
- With our shiny new `is_nothrow_swappable` and the above metafunctions, we can determine if our swap is noexcept.
 - We only emit the swap if it is noexcept.
 - Assignment operators won't be available otherwise.

Adding member and ADL swap

- [error_or.hpp: L106-L112](#)
 - Our member swap delegates to our SFINAE'd `sfinae_swap`.
 - Our ADL swap delegates to member swap.
 - Both members expose `noexcept` status.
 - `detail::is_swappable` works on `error_or<T>`.
 - `detail::is_nothrow_swappable` works on `error_or<T>`.

Move and noexcept are friends

- If you must write a move special member despite the rule of zero, and you can prove that it doesn't throw, then you should mark it with `noexcept`.
 - Allows templates instantiated on your type to determine that moves can be used safely.
 - Allows templates instantiated on your type to propagate `noexcept` on their own special member functions.
 - Enables non-throwing [`std::swap`](#) or ADL swap via `std::move`.
 - Which enables exception safe assignment via copy and swap.

Accessors: Some easy stuff...

- [error_or.hpp: L143-L174](#)
 - ‘ok()’ tells us whether we have a value or an error.
 - operator bool() for use in conditionals
 - value and error accessors return the identified object.
 - ‘release’ functions to move values or errors out.

Integration with `<system_error>`

- [error_or.hpp: L243-L247](#)
 - Use C++11 ‘template typedefs’ via ‘using’ to define nice names for `error_or<std::error_[condition|code], T>`
 - No need to roll your own ‘error type’, just make a new category and use what the standard gives you.
 - These types are noexcept.

Integration with [std::unique_ptr](#)

- [error_or.hpp: L249-L253](#)
 - `std::unique_ptr` makes a great return value for functions that create a new object, and has noexcept move semantics.
 - Combining it with `error_code_or` gives you the following nice semantics:
 - The function returns to me an owned pointer to a newly created object, or an error explaining why it was not constructed.
 - Example:
 - `error_code_or_unique<socket> connect_to(std::string host);`

Fun and Games

- Convert a function that returns `T` or throws a `std::system_error` into a function that returns `error_code_or<T>`:
 - [converters.hpp: L25-L45](#)
 - I suspect there is a better way to declare this, but I haven't derived it yet.
- Convert a function that returns `error_code_or<T>` into one that returns `T` or throws:
 - [converters.hpp: L47-L65](#)
 - Again, there might be a better way to do this.

Fun and Games: In Progress

- I'd like this to work:

```
error_code_or_unique<socket> connect(std::string host) {...}
error_code_or_unique<request> socket::send(std::string message)
{...}

auto response = connect("www.10gen.com").switch(
    [] (std::unique_ptr<socket>&& sock) ->
    error_code_or_unique<response> {
        return sock->send("GET /index.html http/1.1").switch(
            [sock] (std::unique_ptr<request>&& request) {
                return request->await_response(std::move(sock));
            },
            [] (std::error_code&& error) ->
            error_code_or_unique<response>{
                std::cerr << "Failed to send: " << error << std::endl;
                return error;
            }
        )
    },
    [] (std::error_code&& error) -> error_code_or_unique<response> {
        std::cerr << "Failed to connect: " << error << std::endl;
        return error;
    }
);
```


Examples and Demo

Thanks!