

Introduction to C++ smart pointers

Jason Rassi

rassi@mongodb.com

github.com/jrassi



Smart pointer pattern

- Using pointer-like objects to make programs simple and leak-free

Resource leaks

```
f (g (), new Kitten);
```

Assumptions

- This is a beginner talk
- You're familiar with use of basic C++ language features
 - Classes, templates, pointers, references, exceptions, operator overloading
- You don't know what a smart pointer is, or are curious to learn more about them
- You like kittens

Roadmap

- Pointers and ownership
- What is a smart pointer?
- Survey
 - `boost::scoped_ptr`
 - `std::auto_ptr`
 - `std::unique_ptr`
 - `std::shared_ptr`, `std::weak_ptr`

What is a pointer?


- Conceptually: a handle to an object
 - Implements indirection
 - Often can be re-bound to a different object
- Pointers (“bare pointers”) in C/C++:
 - Associated with a specific type named in declaration
 - Generally are in one of two valid states
 - Point to an object
 - Value is the starting location of that object in the process’ address space
 - Point to no objects
 - Value is null
 - Simple data type: no destructor

Kittens

- Properties of a Kitten
 - Lifetime is independent of lifetime of other objects in world
 - At any point in time, all Kittens must have [at least] one owner
 - Kitten with no owner is a leaked Kitten
 - Billions of Kittens with no owner => obvious system stability issues

Pointers in action

```
Kitten k_obj;  
Kitten* k1 = &k_obj;  
Kitten* k2 = new Kitten;  
Kitten* k3 = nullptr;
```



What is ownership?

- Loose definition: responsibility for cleaning up after
- How do owners of these resources “clean up?”
 - Dynamically-allocated memory
 - File descriptors
 - Network connections
 - Locks

Ownership

- I found a pointer. Should I call delete on it?
 - Certainly not if the associated object has static or automatic lifetime
 - Certainly not if it's already been deleted
 - Certainly not if it's someone else's job to delete it
 - Certainly not if the associated object was allocated with a custom allocator

```
Kitten* k = a_mystery_function();
```

Ownership policies

- Can an owner give away the object?
 - Transferrable ownership
 - Non-transferrable ownership
- Can the object have multiple owners at once?
 - Shared ownership
 - Exclusive ownership

Exclusive, non-transferrable ownership

- Once owner, forever owner
 - Person acquires Kitten
 - Person never lets anyone else borrow or share Kitten
 - Kitten's lifetime must not outlast Person's lifetime

Exclusive, transferrable ownership

- One owner at a time, but owner can change
 - Person acquires Kitten
 - Person can use kitten-sitter for vacation:
 - Before vacation, person transfers ownership to kitten-sitter
 - After vacation, kitten-sitter transfers ownership back
 - Person can give Kitten up for adoption
 - Kitten's lifetime must not outlast its current owner (but can outlast previous owners)

Shareable ownership

- Set of owners changes over time
 - Person acquires Kitten
 - Ownership of Kitten shared with cohabitants, descendants
 - Owners added, removed over time
 - Kitten's lifetime must not outlast that of its last remaining owner

Ownership with raw pointers

- Error-prone!
- Ownership transfer

```
find_new_owner(k);  
k = nullptr;  
// k not valid
```

- Ownership sharing

```
add_new_roommate(candidates["alice"], k);  
// k still valid
```

Pitfalls of bare pointers

- Suppose we're play-testing "t", a new Kitten toy
 - Kitten promises not to throw in `play_with_toy()`

```
{
    Kitten* k = kitten_factory();
    Rating r = k->play_with_toy(t);
    delete k;
    results.add_rating(r);
}
```
- Simple to write first pass, but simple to maintain?

Pitfalls of bare pointers

- What if the precondition on `play_with_toy()` changes?
- What if `play_with_toy()`'s exception specification changes?

```
{
    Kitten* k = kitten_factory();
    if (!k->interested_in_play()) {
        delete k;
        return;
    }
    Rating r;
    try {
        r = k->play_with_toy(t);
    }
    catch (CutenessException& e) {
        delete k;
        throw;
    }
    delete k;
    results.add_rating(r);
}
```

What happened?

- We're lazy
 - Putting “delete k” in the exact correct set of places is hard work
- What all those “delete k” statements are trying to express here:
 - We just want the Kitten object to be deleted before it goes out of scope

Roadmap

- ~~Pointers and ownership~~
- What is a smart pointer?
- Survey
 - `boost::scoped_ptr`
 - `std::auto_ptr`
 - `std::unique_ptr`
 - `std::shared_ptr`, `std::weak_ptr`

What is a “smart pointer?”

- Loose definition: object that behaves like a pointer, but somehow “smarter”
 - Major similarities to bare pointer:
 - Is bound to 0 or 1 objects at a time, often can be rebound
 - Supports indirection: operator*, operator->
 - Major differences:
 - Has some “smart feature”

“Smart feature”

- Traditional smart pointers: resource management
 - Automatic deletion of the owned object
- Other “smart pointers”:
 - Iterators (think: `*it`, `it->f()`, `it++`)
 - Objects with pointer semantics that do something silly when dereferenced

“Smart feature”

- Traditional smart pointers: resource management
 - Automatic deletion of the owned object
- Other “smart pointers”:
 - Iterators (think: `*it`, `it->f()`, `it++`)
 - Objects with pointer semantics that do something silly when dereferenced

```
$ cowsay -f hellokitty 'Meow!'
  _____
< Meow! >
  -----
  \
  /
  ^__^
  (oo)\_____
  | 0 . 0 |
  |-----|
$
```

```
template<typename T>
class kittensay_ptr {
    T* ptr_;

public:
    explicit kittensay_ptr(T* t) : ptr_(t) {}

    T& operator*() {
        std::cout << "Meow!\n";
        return *ptr_;
    }

    T* operator->() {
        std::cout << "Meow!\n";
        return ptr_;
    }
};
```

kittensay_ptr

- In action:

```
int main() {  
    Kitten k_obj;  
    kittensay_ptr<Kitten> k(&k_obj);  
    k->feed();  
    return EXIT_SUCCESS;  
}
```

```
$ ./a.out
```

Meow!

operator->()

- Unary operator
 - Even though it has a thing on the left and on the right
- Evaluating “f->method()” – pseudocode:
 - $x := f$
 - while x not a bare pointer:
 - $x := x.operator->()$
 - evaluate x->method()
- Temporary objects created in while loop

Resource Acquisition Is Initialization

- RAII is the pattern of:
 - Acquiring resources with a constructor
 - Releasing resources in a destructor
- Why?
 - Guarantees that the acquired resource is not leaked
 - Simpler code
- Smart pointer pattern is an application of the general RAII pattern

Smart pointers to the rescue

- Suppose we wrapped bare Kitten pointers in a MyKittenPtr object:

- `MyKittenPtr::MyKittenPtr(Kitten* k) : k_(k) {}`
- `Kitten* MyKittenPtr::operator->() { return k_; }`
- `MyKittenPtr::~~MyKittenPtr() { delete k_; }`

```
{  
    MyKittenPtr k(kitten_factory());  
    if (k->interested_in_play()) {  
        Rating r = k->play_with_toy(t);  
        results.add_rating(r);  
    }  
}
```

Smart pointers to the rescue

- Idea: templatize MyKittenPtr as MyPtr
 - `MyPtr<T>::MyPtr(T* p) : p_(p) {}`
 - `T* MyPtr<T>::operator->() { return p_; }`
 - `MyPtr<T>::~~MyPtr() { delete p_; }`
- ```
{
 MyPtr<Kitten> k(kitten_factory());
 if (k->interested_in_play()) {
 Rating r = k->play_with_toy(t);
 results.add_rating(r);
 }
}
```
- MyPtr is now reusable

# Standard smart pointers

- RAII class
- Stores a pointer to an object on the free store
  - Smart pointer owns the object
    - Hence responsible for deleting it
- Pointer semantics
- Self-documents the object's ownership policy

# Roadmap

- ~~Pointers and ownership~~
- ~~What is a smart pointer?~~
- Survey
  - `boost::scoped_ptr`
  - `std::auto_ptr`
  - `std::unique_ptr`
  - `std::shared_ptr`, `std::weak_ptr`

# boost::scoped\_ptr

- Exclusive, non-transferrable ownership
  - Not copyable, not moveable
  - Once owner, forever owner
- Available in Boost, not C++ standard library
  - One of the first smart pointer classes in Boost
- How it gets its name:
  - Owned object guaranteed\* to be deleted by the time `scoped_ptr` goes out of scope
    - \*Unless `swap()` called – supporting `swap()` is arguably a design error

# Common smart pointer methods

- `void sp::reset(T* k = nullptr);`
  - Replace my existing Kitten with a new Kitten
  - My former Kitten is deleted
- `T* sp::get() const;`
  - Expose to the caller a bare pointer to my Kitten, without changing ownership
    - For use with a legacy API only



```

template<typename T>
class scoped_ptr {
 T* px_;

 scoped_ptr(const scoped_ptr&);
 scoped_ptr& operator=(const scoped_ptr&);

public:
 explicit scoped_ptr(T* p = 0): px_(p) {}

 ~scoped_ptr() { delete px_; }

 T& operator*() const { return *px_; }

 T* operator->() const { return px_; }

 T* get() const { return px_; }

 void swap(scoped_ptr& b) {
 T* tmp = b.px_;
 b.px_ = px_;
 px_ = tmp;
 }

 void reset(T* p = 0) { scoped_ptr<T>(p).swap(*this); }

};

```

# boost::scoped\_ptr

- Example

```
{
 boost::scoped_ptr<Kitten> a(new Kitten);
 a->feed();
 boost::scoped_ptr<Kitten> b;
 b = a; // ERROR: will not compile.
 b.reset(new Kitten);
 b->feed();
 b.reset(new Kitten);
 b->feed();
}
```

# boost::scoped\_ptr

- One object only
  - Storing a pointer to a C-style array of objects: undefined behavior
  - If you must store an array, use boost::scoped\_array
- No space overhead, almost no runtime overhead
  - `sizeof(boost::scoped_ptr<T>) == sizeof(T*)`
  - Inlined methods
    - Expands to nearly the same code as what you'd write with a bare pointer

# std::auto\_ptr

- Exclusive, transferrable ownership
- Introduced in C++98 as first standard-defined resource management smart pointer
- One object only
  - Incorrect to store a pointer to a C-style array in an auto\_ptr
- Deprecated in C++11 in favor of std::unique\_ptr

# Common smart pointer methods

- `T* sp::release();`
  - Release ownership of my Kitten to the caller
  - Why not in `boost::scoped_ptr`?
    - Because `boost::scoped_ptr` provides non-transferrable ownership semantics

```
{
 std::auto_ptr<Kitten> a(new Kitten);
 std::auto_ptr<Kitten> b(a.release());
 // b now owns the Kitten.
 a->feed(); // ERROR: null dereference.
 b->feed(); // Kitten is fed.

 a.reset(b.release());
 // a owns the Kitten again.
 a->feed(); // Kitten is fed.
 b->feed(); // ERROR: null dereference.
}
```

# std::auto\_ptr

- What else is different from boost::scoped\_ptr?
  - More ways to transfer: “copy” constructor and “copy” assignment operator
    - They take a non-const reference (!)
- What?

```
{
 std::auto_ptr<Kitten> a(new Kitten);
 std::auto_ptr<Kitten> b;
 b = a; // b now owns the Kitten.
 a->feed(); // ERROR: null dereference.
 b->feed(); // Kitten is fed.
}
```

# Move versus copy

- Copy A to B: deep clone
  - Typically: member-wise deep copy
- Move A to B: rip the guts out of A and stuff them in B
  - Typically:
    - Member-wise shallow copy
    - Reset all of A's members to default values

```
template<typename T>
class auto_ptr {
 T* px_;

public:
 explicit auto_ptr(T* p = 0) : px_(p) { }

 ~auto_ptr() { delete px_; }

 T& operator*() const { return *px_; }

 T* operator->() const { return px_; }

 T* get() const { return px_; }

 ...
}
```



...

```
T* release() {
 T* tmp = px_
 px_ = 0;
 return tmp;
}
```

```
void reset(T* p = 0) {
 if (p != px_) {
 delete px_
 px_ = p;
 }
}
```

```
auto_ptr(auto_ptr& a) : px_(a.release()) { }
```

```
auto_ptr& operator=(auto_ptr& a) {
 reset(a.release());
 return *this;
}
```

```
};
```

# std::auto\_ptr

- You will be in hot water if you pass auto\_ptr to a template function that executes “a = b;” and expects a copy
  - Compiler will make a valiant effort to prevent this by enforcing const-correctness
    - Only if developer didn’t forget to const-qualify...
- Can’t be put in standard containers

# std::auto\_ptr

- If:
  - You are able to use a compiler with C++11 support in your build system
- Then:
  - Do not use std::auto\_ptr
  - Instead, use std::unique\_ptr

# `std::unique_ptr`

- Exclusive, transferrable ownership
- New in C++11
  - What `std::auto_ptr` always wanted to be, but the language didn't support

# std::unique\_ptr

- Can store a single pointer or a C-style array
  - Single pointer: `unique_ptr<T>`
  - Array: `unique_ptr<T[]>`
- Custom deleter: world of possibilities!
  - Single pointer: `unique_ptr<T, D>`
  - Array: `unique_ptr<T[], D>`
- So is `unique_ptr`'s “move” any better than `auto_ptr`'s?

# Rvalue references lightning talk

- In C++11, rvalue reference (“T&&”) introduced
  - Similar to an lvalue reference (“T&”)
  - But different rules for binding to expressions:
    - Can be bound to a temporary
    - Can’t be bound to a “normal” lvalue without an explicit cast
- Gives rise to the following idiom
  - `f(T&& k) { ... }`
    - Caller of `f` is saying “Here, take `k` and do what you wish with it. I promise never to look at it again.”
    - “Move constructor”, “move assignment operator”
- Why relevant to this talk? Ownership transfer!

# Rvalue references lightning talk

- `std::move(x)`
  - Poorly-named
  - Performs a cast
    - Casts argument to an rvalue reference
  - Does not perform a move
    - Generates zero code
- Move constructor and move assignment operator take a “T&&”
  - Hence the idiom “`a = std::move(b);`”
    - Overload resolution makes the compiler generate a call to “`operator=(T&& t);`”

# std::unique\_ptr

- Move: allowed
  - Move constructor:
    - `unique_ptr(unique_ptr&& u);`
  - Move assignment operator
    - `unique_ptr& operator=(unique_ptr&& u);`
- Copy: not allowed
  - No copy constructor:
    - `unique_ptr(const unique_ptr&);` // Not defined.
  - No copy assignment operator:
    - `unique_ptr& operator=(const unique_ptr&);` // Not defined.



```

/** Simplification: no custom deleter, no support for arrays. */
template<typename T>
class my_unique_ptr {
 T* px_;

public:
 explicit my_unique_ptr(T* p = 0) : px_(p) { }

 ~my_unique_ptr() { delete px_; }

 T& operator*() const { return *px_; }

 T* operator->() const { return px_; }

 T* get() const { return px_; }

 T* release() {
 T* tmp = px_;
 px_ = 0;
 return tmp;
 }

 void reset(T* p = 0) {
 if (p != px_) {
 delete px_;
 px_ = p;
 }
 }
}

```

...

...

public:

```
 my_unique_ptr(my_unique_ptr&& u)
 : px_(u.release()) { }
```

```
 my_unique_ptr& operator=(my_unique_ptr&& u) {
 reset(u.release());
 return *this;
 }
```

private:

```
 my_unique_ptr(const my_unique_ptr&);
 my_unique_ptr& operator=(const my_unique_ptr&);
```

```
};
```

# std::unique\_ptr

- What does assignment look like?

```
{
 std::unique_ptr<Kitten> a(new Kitten);
 std::unique_ptr<Kitten> b;
 b = a; // ERROR: will not compile.
 b = std::move(a); // b now owns the Kitten.
 a->feed(); // ERROR: null dereference.
 b->feed(); // Kitten is fed.
}
```

# std::unique\_ptr

- Good fit for a factory function

```
std::unique_ptr<Kitten> kitten_factory();
```

- Good fit for standard containers

```
std::vector<std::unique_ptr<Kitten>> v;
```

- Good fit for instance variable (if pointer semantics required)
  - No special handling needed for moveable classes
  - Special handling needed for copyable classes
- And most places where you would have a raw pointer to an object with dynamic lifetime
  - Except when shared ownership is absolutely necessary

# std::unique\_ptr custom deleter

- Pass a function object or function pointer as the second argument to the std::unique\_ptr() ctor
- Allows storing of pointers to objects that need special freeing logic
  - Perhaps you're using a pool allocator
- Where is the deleter stored?
  - As additional private member data
  - In practice:
    - sizeof(std::unique\_ptr<T, D>) > sizeof(T\*)
    - sizeof(std::unique\_ptr<T>) == sizeof(T\*)

# std::make\_unique()

- Performs the allocation for you using new
- Accepted into C++14
- Never write a bare call to new again!

```
std::unique_ptr<Kitten> k(std::make_unique<Kitten>());
```

# std::make\_unique()

- What's wrong with the bare “new” call?
- Consider:

```
void f(bool b, std::unique_ptr<Kitten> k);
```

- Is it possible for the Kitten to be leaked here?

```
f(g(), std::unique_ptr<Kitten>(new Kitten));
```

- What about here?

```
f(g(), std::make_unique<Kitten>());
```

# std::shared\_ptr

- Shared ownership
  - Thread-safe
- History
  - Started in Boost
    - boost::shared\_ptr
  - Pulled into TR1
    - std::tr1::shared\_ptr
  - Pulled into std namespace for C++11
    - std::shared\_ptr



# std::shared\_ptr


- Supports custom deleter
- Single object only
  - Use boost::shared\_ptr (1.53 or newer) for arrays
    - boost::shared\_ptr<T[]>, or
    - boost::shared\_ptr<T[N]>
  - Older versions of boost: use boost::shared\_array<T>

# std::shared\_ptr

- Reference-counting smart pointer
  - Not a garbage collector
- Reference count stored in dynamically-allocated *control block*
  - Control block allocated in “first owner” constructor, deallocated in “last owner” destructor
  - Reference count updated when a new owner is added or removed
    - Atomic increment/decrement instructions

# std::shared\_ptr control block

```
std::shared_ptr<Kitten> s1(new Kitten);
std::shared_ptr<Kitten> s2(s1);
```



# std::shared\_ptr

- When is associated object deleted?
  - shared\_ptr destructor checks reference count
  - If no more owners, deletes object


```
std::shared_ptr<Kitten> o1(new Kitten);
std::shared_ptr<Kitten> o2;
o2 = o1; // Owners now: o1, o2.
o1 = nullptr; // Owners now: o2.
o1 = std::move(o2); // Owners now: o1.
o1 = nullptr; // No owners! Deletion now.
```

# std::shared\_ptr, std::weak\_ptr


- shared\_ptr to object: “strong owner” of object
- weak\_ptr to same object: “weak owner”
  - Think “observer from afar”
  - Can’t access object directly: no operator->() etc.
- Control block contains “strong reference count” and “weak reference count”
  - Object deleted when strong count reaches zero
    - Even if weak owners still exist
  - Control block deleted when strong count and weak count reach zero
- Use weak\_ptr for:
  - Statistics tracking
  - Breaking cycles of strong owners

# Control block, now with std::weak\_ptr

```
std::shared_ptr<Kitten> s1(new Kitten);
std::weak_ptr<Kitten> w1(s1);
```




# Ownership cycle: tree with shared\_ptr back references




# Ownership cycle: tree with shared\_ptr back references

```
this.right_child_.reset(); // Sub-tree is leaked.
```






# Breaking the cycle: tree with weak\_ptr back references



# Breaking the cycle: tree with weak\_ptr back references

```
this.right_child_.reset(); // Sub-tree is deleted.
```



# std::weak\_ptr

- lock()
  - `shared_ptr<T> weak_ptr<T>::lock() const;`
  - Weak owners use this to request strong ownership
  - If object expired, returns empty `shared_ptr<T>()`
  - Otherwise, returns `shared_ptr<T>(*this)`

# std::shared\_ptr, std::weak\_ptr

- In action:


```
std::shared_ptr<Kitten> s1(new Kitten);
std::shared_ptr<Kitten> s2;
std::weak_ptr<Kitten> w;
w = s1; // Owners now: s1, w.
s2 = w.lock(); // Owners now: s1, s2, w.
s1 = s2 = nullptr; // Owners now: w. No strong owners!
 // Kitten deletion now.
s2 = w.lock(); // Owners now: w. s2 does not
 // change.
w = nullptr; // No owners at all! Control block
 // deallocation.
```

# std::make\_shared()

- std::make\_unique() for shared pointers
  - Unlike std::make\_unique(), this is actually in C++11
- std::make\_shared() allocates the control block and the managed object in a single allocation
  - Pro: one allocation instead of two
  - Con: memory associated with object won't be freed until all owners (strong and weak) are gone

# Control block, now with `std::make_shared()`

```
std::shared_ptr<Kitten> s1(std::make_shared<Kitten>());
std::weak_ptr<Kitten> w1(s1);
s1 = nullptr; // Kitten destroyed, memory still alloc.
```



# Be wary of shared ownership

- “Do not design your code to use shared ownership without a very good reason”
  - “One such reason is to avoid expensive copy operations, but you should only do this if the performance benefits are significant, and the underlying object is immutable (i.e. `shared_ptr<const Foo>`)”
- “If you do use shared ownership, prefer to use `shared_ptr`”

# Review

|                         | Ownership | Copyable?<br>Moveable? | Availability      |
|-------------------------|-----------|------------------------|-------------------|
| scoped_ptr              | Exclusive | --                     | Boost             |
| auto_ptr                | Exclusive | Moveable... poorly     | C++98             |
| unique_ptr              | Exclusive | Moveable               | C++11             |
| shared_ptr,<br>weak_ptr | Shared    | Copyable<br>Moveable   | Boost, TR1, C++11 |



# boost/smart\_ptr/

The smart pointer library provides six smart pointer class templates:

|                                      |                                                        |                                                                |
|--------------------------------------|--------------------------------------------------------|----------------------------------------------------------------|
| <a href="#"><u>scoped_ptr</u></a>    | <a href="#"><u>&lt;boost/scoped_ptr.hpp&gt;</u></a>    | Simple sole ownership of single objects. Noncopyable.          |
| <a href="#"><u>scoped_array</u></a>  | <a href="#"><u>&lt;boost/scoped_array.hpp&gt;</u></a>  | Simple sole ownership of arrays. Noncopyable.                  |
| <a href="#"><u>shared_ptr</u></a>    | <a href="#"><u>&lt;boost/shared_ptr.hpp&gt;</u></a>    | Object ownership shared among multiple pointers.               |
| <a href="#"><u>shared_array</u></a>  | <a href="#"><u>&lt;boost/shared_array.hpp&gt;</u></a>  | Array ownership shared among multiple pointers.                |
| <a href="#"><u>weak_ptr</u></a>      | <a href="#"><u>&lt;boost/weak_ptr.hpp&gt;</u></a>      | Non-owning observers of an object owned by <b>shared_ptr</b> . |
| <a href="#"><u>intrusive_ptr</u></a> | <a href="#"><u>&lt;boost/intrusive_ptr.hpp&gt;</u></a> | Shared ownership of objects with an embedded reference count.  |

# Further reading

- Meyers, Scott. *Effective C++, Third Edition*.
- Stroustrup, Bjarne. *The C++ Programming Language, Fourth Edition*.
- Sutter, Herb. *More Exceptional C++*.

# Special thanks

- Andrew Morrow
- Charlie Page
- Greg Steinbruner

# Resource leaks no more

```
f (g (),
 std::make_unique<Kitten> ()) ;
```